

DeWang Li(Nathon Lee)

leejianwoo@gmail.com — github.com/nathon-lee — +86 191-4138-0850

Summary

AI Infrastructure and Distributed Systems Engineer with experience in large-scale training systems, high-performance storage, and low-latency services. Active contributor to DeepSpeed and open-source LLM infrastructure projects. Experienced with PyTorch, DeepSpeed, Megatron-LM, RLHF, PPO, DPO, and distributed training technologies including DP, TP, PP, CP, EP, FSDP, and ZeRO.

Open Source Contributions

DeepSpeed Contributor (Microsoft)

- Contributed multiple production-grade pull requests to DeepSpeed.
- Participated in implementing **AutoSP (Automatic Sequence Parallelism)** support for multimodal training workloads.
- Enabled Sequence Parallelism for ViT + LLM architectures to improve scalability and memory efficiency.
- Implemented automatic best-checkpoint saving based on evaluation metrics.
- Improved training stability and framework usability.

Selected Contributions

DeepSpeed: #7984, #7908, #7928, #7763, #7940, #7934 (AutoSP Contributor)

OpenRLHF #1206 — **RLlib** #61823 — **Tinker** #172, #286

Technical Skills

Languages: Python, C/C++, Go, Java, Bash

Training: PyTorch, DeepSpeed, Megatron-LM, RLHF, PPO, DPO, LoRA

Parallelism: DP, TP, PP, CP, EP, FSDP, ZeRO, FP8

Systems: Distributed Systems, NUMA, RPC, High-Performance Storage, Performance Optimization

Professional Experience

JD.com

2025 – Present

Senior Distributed Systems Engineer

- Developed and optimized a large-scale image infrastructure platform processing over **1B+ images daily** and storing more than **300B+ images**.
- Replaced Tomcat → Nginx HTTP service chains with direct JNI-based invocation.
- Eliminated cross-process communication and serialization overhead.
- Achieved **40% reduction in TP99 latency** and **35% reduction in average latency**.

Huawei

2022 – 2024

Senior Distributed Systems Engineer (AI Storage)

- Optimized distributed storage systems for AI training workloads and high-throughput data access.
- Introduced metadata prefetching optimization in the storage read path.
- Increased throughput from **120K OPS** to **180K OPS (+50%)**.
- Designed NUMA-aware scheduling and CPU affinity strategies.
- Improved single-node bandwidth by **5 GB/s**.

Education

University of Electronic Science and Technology of China

2019 – 2022

M.S. in Computer Science

Project

Edge AI System

- Built a distributed edge AI platform using Kubernetes and KubeEdge for resource-constrained AI deployment.